

Niranjan Kumar Gurram

Arlington, TX • +1 (737) 270-7121 • niranjangurram@gmail.com • [linkedin.com/in/niranjang07](https://www.linkedin.com/in/niranjang07)

SUMMARY

Results-driven data engineer blending IaC, streaming, and MLOps to deliver audit-ready, low-latency pipelines. Proven record cutting MTTR, cloud spend, and churn while raising NPS and analyst capacity in finance & telecom. Skilled at leading cross-functional squads, mentoring juniors, and translating technical gains into \$-level impact. Currently seeking full-time roles to scale event-driven, governed data platforms that power high-stakes decisions.

SKILLS

PROGRAMMING: Python, SQL, Java, JavaScript, C++

DATA ENGINEERING: Spark, Kafka, Airflow, Debezium, dbt

CLOUD & DEVOPS: AWS, Azure, GCP, Terraform, Kubernetes, ArgoCD, Helm

DATA GOVERNANCE: Azure Purview, RBAC, Great Expectations, GDPR/SOX

ANALYTICS & ML: PyTorch, TensorFlow, Scikit-learn, Power BI, Tableau

PROFESSIONAL EXPERIENCE

Graduate Data Engineer Intern

Capital One

Feb 2024 - Jun 2025

- Orchestrated 38-module Terraform stack across AWS, Azure, GCP; drift checks cut IaC tickets 45% and saved \$260k.
- Built Kafka-Debezium CDC pipeline moving 11k TPS events to Snowflake in sub-sec SLA, enabling real-time fraud analytics.
- Productionized PyTorch fraud model with AUC 0.94 and 27% false-positive drop; feature store, drift alerts cut triage 60%.
- Implemented Helm charts & ArgoCD GitOps; blue-green rollbacks avg 3 min, cutting MTTR 22→12 min each release.
- Deployed Prometheus-Grafana SLO dashboards; alerts cut incident MTTR 38→22 min and avoided \$80k SLA penalties.
- Managed 220-item backlog; sprint velocity rose 18% and on-time releases lifted analytics team NPS 54→72.
- Mentored 4 junior engineers; onboarding time dropped 9→5 weeks and team retention improved by 12%.

Data Engineering Co-op

Tata Consultancy Services

Jul 2021 - Aug 2023

- Built 20-node Spark cluster processing 2 TB telecom data daily; analyst wait time fell 55%, freeing 360 hours yearly.
- Optimized Spark partitions and Z-order; shuffle spill dropped 47%, KPI refresh halved, saving \$110k compute expense.
- Authored 17 Airflow DAGs with Slack & PagerDuty hooks; MTTR plummeted 120→38 min and SLA breaches shrank 20%→2%.
- Automated KPI pipelines with Python/SQL and dbt; SLA compliance climbed 80→98%, unlocking 14 h monthly capacity.
- Refactored Tableau views to extracts; load 18→12 s and adoption up 32%, influencing \$4.7 M retention decisions.
- Rolled out Azure Purview: 86 RBAC policies, 210 masked columns, lineage jobs; cleared two GDPR audits with zero issues.
- Embedded Great Expectations at 85% coverage plus JMeter load gates; defect density fell 42% and avoided 6 hotfixes.

PROJECTS

Facial Insight

- Flask + TensorFlow attendance system recognizing 50 faces/frame at 96% accuracy, auto-logging to MySQL.

Real-Time Fare Estimator

- Kafka → Spark Streaming predicting rideshare cost within 7% MAE on 1.2 M trips, served via REST.

EDUCATION

University of Texas at Arlington - Arlington, TX

M.S. Computer Science & Information Systems

Aug 2023 - May 2025

- Relevant Coursework:* Algorithms, Data Mining, Cloud Computing, Big Data, Machine Learning

Vignan's University - India

B.Tech. Computer Science

Aug 2019 - May 2023

CERTIFICATIONS

-
- AWS Certified Data Analytics – Specialty (In Progress)
 - HashiCorp Terraform Associate (2024)
 - Microsoft Azure Data Engineer DP-203 (2025)